

AG News Classification with Transformer and DNN Architectures

1. Project Overview

This project investigates AG News classification using two neural architectures:

- A baseline **DNN classifier** - A custom **Transformer-based classifier**

We aim to explore differences in architecture behavior, input handling, and performance under MILA's evaluation style, with a focus on generalization, sequence handling, and training dynamics.

2. Dataset

- **AG News:** News headlines + descriptions
 - **Task:** 4-class classification (World, Sports, Business, Sci/Tech)
 - **Inputs:** Headline + Description (concatenated)
 - **Size:** ~120k training samples
 - **Tokenization:** Word-level regex tokenization (`\b\w+\b`)
 - **Padding:**
 - DNN: static padding (`max_seq_len = 100`)
 - Transformer: dynamic padding per batch with `pad_sequence`
-

3. Model Architectures

3.1 DNNClassifier (Baseline)

- **Embedding** → **MeanPool** → **FC1** → **Dropout** → **FC2** → **Logits**
- Dropout=0.3, no batch norm (BN hurt validation)
- F1: ~91% macro

3.2 TransformerClassifier (Scratch)

- Learned positional encodings
 - 2 encoder layers, 4 heads each
 - Mean pooling with masking (not using [CLS])
 - Masked mean pooling handles <pad> tokens
 - Achieved ~97.4% macro F1 on validation
-

4. Data Pipeline

- Built flexible **TextDataset**:
 - Supports static and dynamic padding

- Switches logic based on `model_type` and `config`
 - Implemented `transformer_collate_fn(pad_idx)`:
 - Uses `pad_sequence` for dynamic batch padding
 - Returns `src_key_padding_mask` to mask attention on padded tokens
 - Config-driven design (`config.yaml`)
-

5. Transformer Forward Pass (Explained)

- Input: `[batch_size, seq_len]` token IDs
 - Embedding + Positional Encoding: `[B, L, D]`
 - TransformerEncoder (2 layers \times 4 heads)
 - Masked Mean Pooling:
 - Uses `src_key_padding_mask`
 - Avoids influence of `<pad>` tokens
 - Final classifier: Linear projection `[D \rightarrow num_classes]`
-

6. Attention Explained (Heads, Masking, and Pooling)

- 4 attention heads look at different token relationships
 - Each token attends to others based on learned relevance (scaled dot product)
 - Head outputs are concatenated $\rightarrow [B, L, D]$
 - Projected using `Linear(D, D)` to blend head info
 - Masked mean pooling aggregates final sequence embeddings robustly
-

7. Training Setup

- Optimizer: Adam
 - Scheduler: ReduceLROnPlateau (monitoring macro F1)
 - Gradient clipping: `clip_grad_norm`
 - Early stopping based on F1 (patience = 5)
 - Stratified validation split (10%)
 - Static vs. dynamic padding evaluated
-

8. Key Observations

- Transformer outperformed DNN by 6–7% F1, even without pretraining
- Mean pooling with masking was essential
- Validation sometimes higher than training \rightarrow possibly due to regularization effects

- Training loss may not reflect generalization performance
 - F1 score chosen as main metric due to class imbalance and MILA recommendation
-

9. Future Work

- Hands-on exploration of Q/K/V tensors per head
 - Deeper look at attention maps and interpretability
 - Extend to pretrained models (e.g., T5, BERT, Flan-T5)
 - Multi-task setup: headline-only vs headline+desc
 - Evaluate Byte-BPE vs word-level for robustness
-

10. Appendix

- Full config file
- Selected training logs
- Visualizations of attention and pooling
- Key implementation notes (collate, masking logic, etc.)